Barcelona Supercomputing Center
Centro Nacional de Supercomputación

EXCELENCIA SEVERO OCHOA

Understanding applications with Paraver

Judit Gimenez
judit@bsc.es

Jan 31st 2019

CRHPCS19, San José

Extrae and Paraver Hands-on

Barcelona Supercomputing Center
Centro Nacional de Supercomputación

Install Paraver

- Download from https://tools.bsc.es/downloads

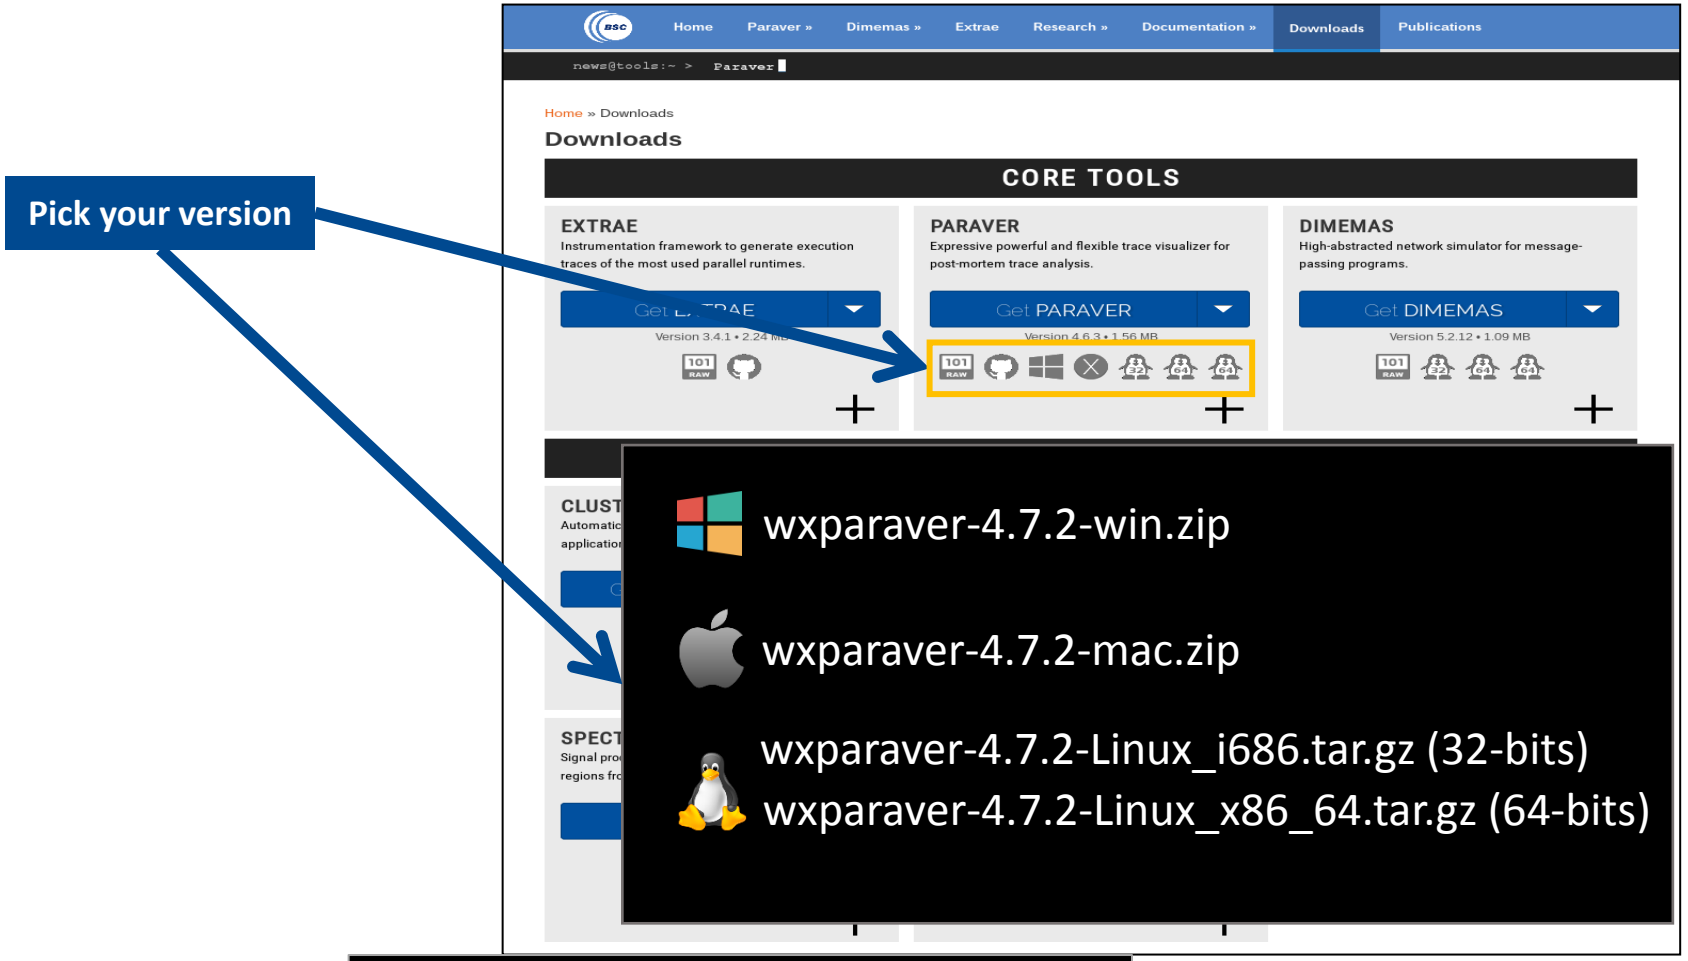

~gimenez1/BSC-paraver

Install Paraver (II)

- Download tutorials:
  - Documentation
    - Paraver tutorials

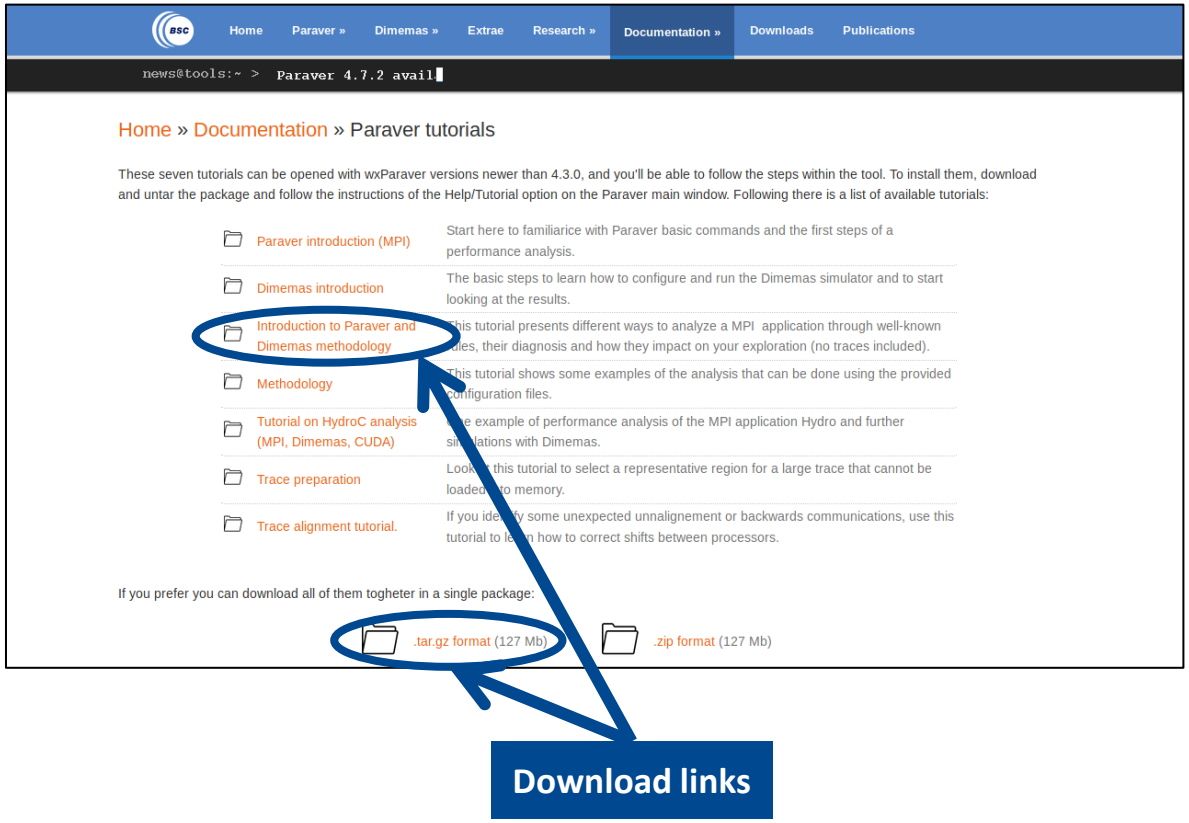

Barcelona Supercomputing Center
Centro Nacional de Supercomputación

~gimenez1/BSC-paraver

Uncompress, rename & move

- Paraver

@ your laptop

```
> tar xf wxparaver-4.8.1-linux-x86_64.tar.gz
> mv wxparaver-4.8.1-linux-x86_64 paraver
```

- Tutorials

@ your laptop

```
> tar xf paraver-tutorials-20150526.tar.gz
> mv paraver-tutorials-20150526 paraver/tutorials
```

Check that everything works

- Start Paraver

@ your laptop

```
> paraver/bin/wxparaver
```

- Check that tutorials are available

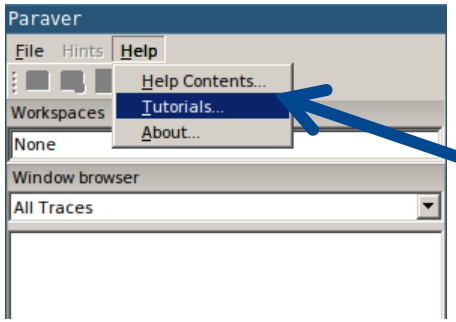

Click on Help → Tutorials

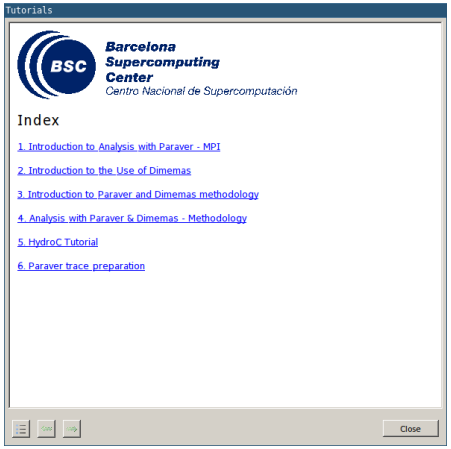

- Trouble installing locally? Remote open from kabre

@ cluster.cenat.ac.cr

```
> ssh -Y <USER>@cluster.cenat.ac.cr
➢ ln –s ~gimenez1/tools/wxparaver64/bin/wxparaver
  wxparaver
```

Log in to kabré

@ your laptop

```
> ssh -Y <USER>@cluster.cenat.ac.cr
```

- Copy the examples to your home folder:

@ cluster.cenat.ac.cr

```
> cp -r ~jgimenez/BSC-handson/ $HOME

> ls -l $HOME/BSC-handson/
    … apps
    … traces
    … slides
```

OpenMP example: matrix

- Comparing the location of the parallel for pragma

```
#pragma omp parallel for
for(int i=0; i<N; i++){
        for(int k=0; k<N; k++){
                for(int j=0; j<N; j++)
```

matrix.l1

```
for(int i=0; i<N; i++){
#pragma omp parallel for
        for(int k=0; k<N; k++){
                for(int j=0; j<N; j++){
```

matrix.l2

- Check the script content
- Submit the job script

@ cluster.cenat.ac.cr

```
> cd $HOME/BSC-handson/apps/matrix
> qsub matrix.pbs
```

OpenMP example: matrix

- Copy the traces to your laptop

@ your laptop

```
> scp <USER>@cluster.cenat.ac.cr: \
    BSC_handson/apps/matrix/matrix.l?.* $HOME
```

- Load the traces with Paraver

@ your laptop

```
> wxparaver matrix.l?.prv*
```

- Trouble getting in the queues? Already available at ~/BSC-handson/traces/matrix/matrix.l?.*

- Compare the two executions with paraver

MPI example: jacobi

- Submit the job script

@ cluster.cenat.ac.cr

```
> cd $HOME/BSC-handson/apps/jacobi
> qsub job.kabre
```

- Copy the traces to your laptop and load the trace with Paraver

@ your laptop

```
> scp <USER>@cluster.cenat.ac.cr: \
    BSC_handson/apps/jacobi/jacobi.???* $HOME
> wxparaver jacobi.prv*
```

- Trouble getting in the queues? Already available at ~/BSC-handson/traces/jacobi/jacobi.???*
- Look at the execution with paraver

Extrae features

- Platforms
  - Intel, Cray, BlueGene, MIC, ARM, Android, Fujitsu Sparc…
- Parallel programming models
  - MPI, OpenMP, pthreads, OmpSs, CUDA, OpenCL, Java, Python…
- Performance Counters
  - Using PAPI interface
- Link to source code
  - Callstack at MPI routines
  - OpenMP outlined routines
  - Selected user functions (Dyninst)
- Periodic sampling
- User events (Extrae API)

No need to recompile / relink!

Extrae overheads

	Average values	Kábre (login)	Kabré (Zárate)
Event	150-200 ns	180 ns	600 ns
Event + PAPI	750 ns – 1 us	580 ns	4.7 us
Event + callstack (1 level)	600 ns	750 ns	3.4 us
Event + callstack (6 levels)	2 us	1.7us	8.2 us

How does Extrae work?

- Symbol substitution through LD_PRELOAD
  - Specific libraries for each combination of runtimes
    - MPI
    - OpenMP
    - OpenMP+MPI
    - …

Recommended

- Dynamic instrumentation
  - Based on Dyninst (developed by U.Wisconsin / U.Maryland)
    - Instrumentation in memory
    - Binary rewriting

- Alternatives
  - Compiler instrumentation (-finstrument-functions – Intel, GNU)
  - Static link (i.e., PMPI, Extrae API)

Using Extrae in 3 steps

1. **Adapt** your job submission scripts

2. (Optional) **Tune** the Extrae XML configuration file
   - Examples distributed with Extrae at $EXTRAE_HOME/share/example
3. **Run** it!

- For further reference check the **Extrae User Guide:**
  - https://tools.bsc.es/sites/default/files/documentation/html/extrae/index.html
  - Also distributed with Extrae at $EXTRAE_HOME/share/doc

Step 1: Adapt the job script to load Extrae

@ cluster.cenat.ac.cr

```
> vi $HOME/BSC-handson/apps/lulesh/job.kabre
```

job.kabre

```
#PBS -N extrae
#PBS -q phi-n1h72
#PBS -l nodes=1:ppn=27
#PBS -l walltime=00:20:00

cd $PBS_O_WORKDIR

module load mpich/3.2.1




mpirun  ./lulesh2.0 -i 20 -s 64 -p
```

Step 1: Adapt the job script to load Extrae

@ cluster.cenat.ac.cr

```
> vi $HOME/BSC-handson/apps/lulesh/job.kabre
```

job.kabre

```
#PBS -N extrae
#PBS -q phi-n1h72
#PBS -l nodes=1:ppn=27
#PBS -l walltime=00:20:00

cd $PBS_O_WORKDIR

module load mpich/3.2.1


export TRACE_NAME=lulesh2_27p.prv


mpirun ./extrae/trace.sh  ./lulesh2.0 -i 20 -s
64 -p
```

Step 1: Adapt the job script to load Extrae

@ cluster.cenat.ac.cr

```
> vi $HOME/BSC-handson/apps/lulesh/extrae/trace.sh
```

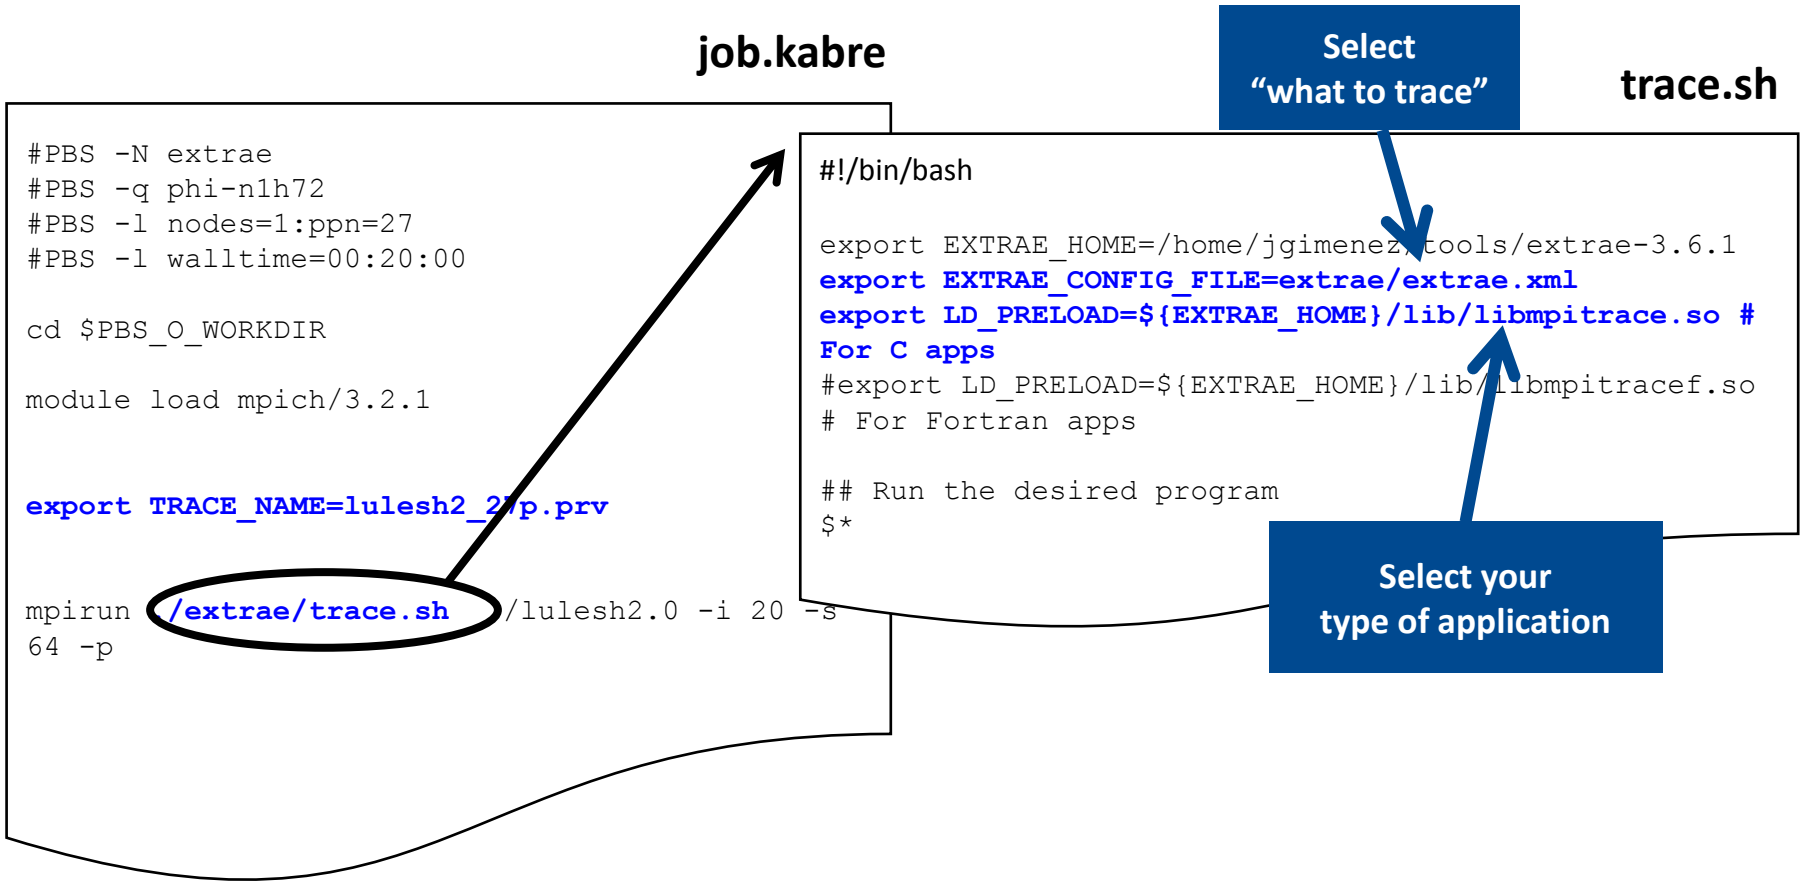

Step 1: Which tracing library?

- Choose depending on the application type

Library	Serial	MPI	OpenMP	pthread	CUDA
libseqtrace	✓				
libmpitrace[f][1]		✓			
libomptrace			✓		
libpttrace				✓	
libcudatrace					✓
libompitrace[f] [1]		✓	✓		
libptmpitrace[f] [1]		✓		✓	
libcudampitrace[f] [1]		✓			✓

[1] include suffix "f" in Fortran codes

Step 3: Run it!

- Submit your job

@ cluster.cenat.ac.cr

```
> cd $HOME/BSC-handson/apps/lulesh

> qsub job.kabre
```

- Easy! ☺

Step 2: Extrae XML configuration

@ cluster.cenat.ac.cr

```
> vi $HOME/BSC-handson/apps/lulesh/extrae/extrae.xml
```

```
<mpi enabled="yes">
  <counters enabled="yes" />
</mpi>

<openmp enabled="no">
  <locks enabled="no" />
  <counters enabled="yes" />
</openmp>

<pthread enabled="no">
  <locks enabled="no" />
  <counters enabled="yes" />
</pthread>

<callers enabled="yes">
  <mpi enabled="yes">1-3</mpi>
  <sampling enabled="no">1-5</sampling>
</callers>
```

Trace the MPI calls
(What's the program doing?)

Trace the call-stack
(Where in my code?)

Step 2: Extrae XML configuration (II)

@ cluster.cenat.ac.cr

```
> vi $HOME/BSC-handson/apps/lulesh/extrae/extrae.xml
```

```
<counters enabled="yes">
  <cpu enabled="yes" starting-set-distribution="1">
    <set enabled="yes" domain="all" changeat-time="500000us">
      PAPI_TOT_INS, PAPI_TOT_CYC, PAPI_L1_DCM, PAPI_L2_DCM, PAPI_BR_MSP, RESOURCE_STALLS
    </set>
    <set enabled="yes" domain="all" changeat-time="500000us">
      PAPI_TOT_INS, PAPI_TOT_CYC, PAPI_L3_TCM, PAPI_LD_INS, PAPI_SR_INS
    </set>
    <set enabled="yes" domain="all" changeat-time="500000us">
      PAPI_TOT_INS, PAPI_TOT_CYC, PAPI_VEC_DP
    </set>
    <set enabled="yes" domain="all" changeat-time="500000us">
      PAPI_TOT_INS, PAPI_TOT_CYC, PAPI_VEC_SP, PAPI_FP_INS
    </set>
  </cpu>
  <network enabled="no" />
  <resource-usage enabled="no" />
  <memory-usage enabled="no" />
</counters>
```

Select which HW counters are measured

Step 2: Extrae XML configuration (III)

@ cluster.cenat.ac.cr

```
> vi $HOME/BSC-handson/apps/lulesh/extrae/extrae.xml
```

```
<buffer enabled="yes">
  <size enabled="yes">5000000</size>
  <circular enabled="no" />
</buffer>

<sampling enabled="no" type="default" period="50m" variability="10m" />

<merge enabled="yes"
  synchronization="default"
  tree-fan-out="16"
  max-memory="512"
  joint-states="yes"
  keep-mpits="yes"
  sort-addresses="yes"
  overwrite="yes"
>
  $TRACE_NAME$
</merge>
```

Trace buffer size (Flush/memory trade-off)

Enable sampling (Want more details?)

Automatic post-processing to generate the trace

All done! Check your resulting trace

- Once finished (check with `squeue`) you will have the trace (3 files):

@ cluster.cenat.ac.cr

```
➢ ls -l $HOME/BSC-handson/apps/lulesh
   ...
   lulesh2_27p.pcf
   lulesh2_27p.prv
   lulesh2_27p.row
```

- Trouble getting in the queues? Already available at ~/BSC-handson/traces/lulesh/lulesh2_27p.*

- Now let's look into it !

First steps of analysis

- Copy the trace to your laptop

@ your laptop

```
> scp <USER>@cluster.cenat.ac.cr: \
    BSC_handson/apps/lulesh/lulesh2_27p.* $HOME
```

- Load the trace with Paraver

Click on File → Load Trace → Browse to "lulesh2.0_27p.prv"

- Follow Tutorial #3

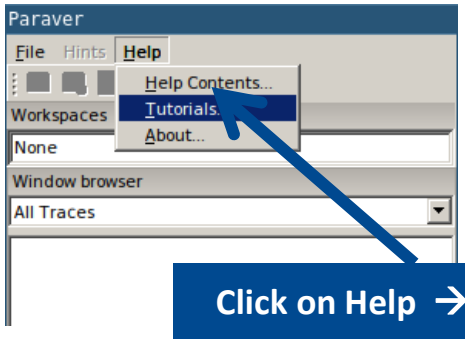

Click on Help → Tutorials

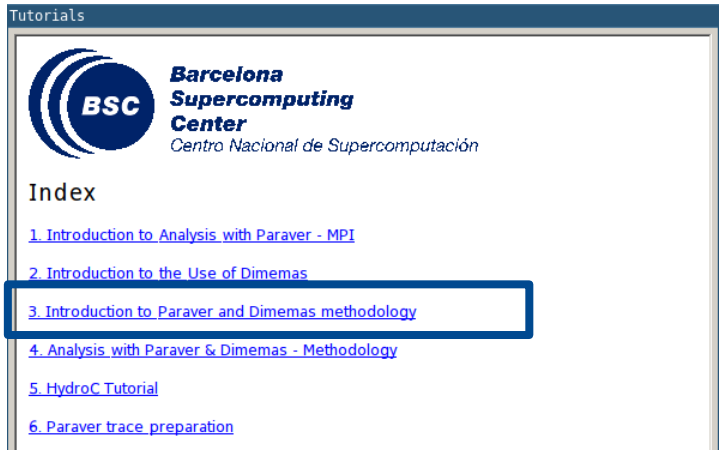

Measure the parallel efficiency

- Click on “mpi_stats.cfg”
  - Check the **Average** for the column labeled “**Outside MPI**”

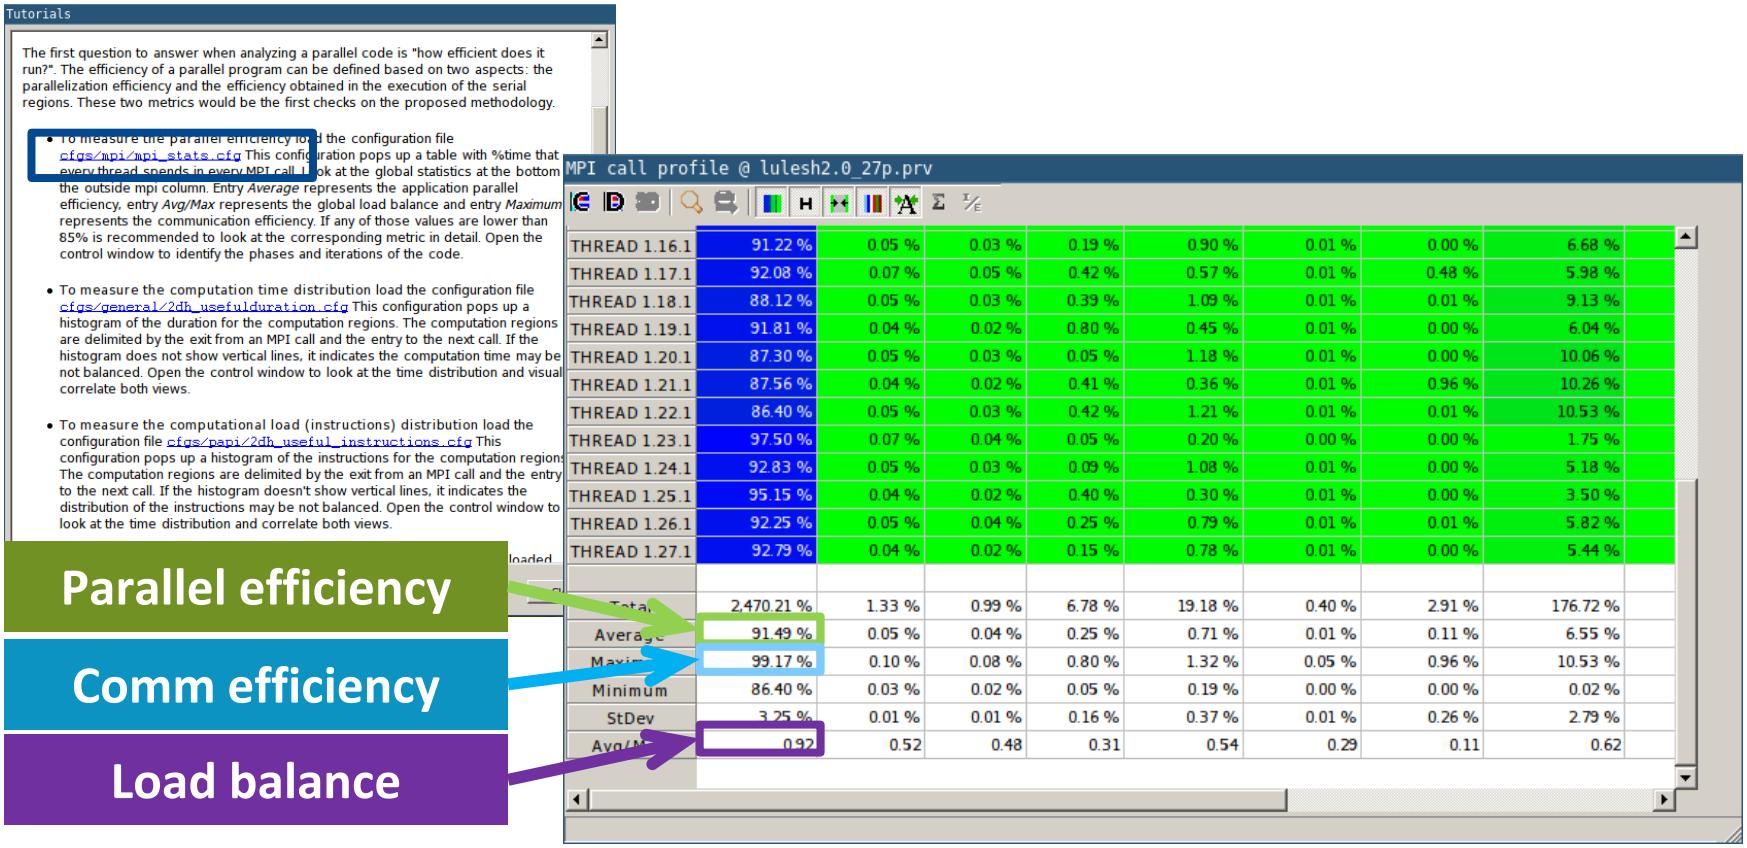

Hints: a good place to start!

- Paraver suggests CFG's based on the information present in the trace

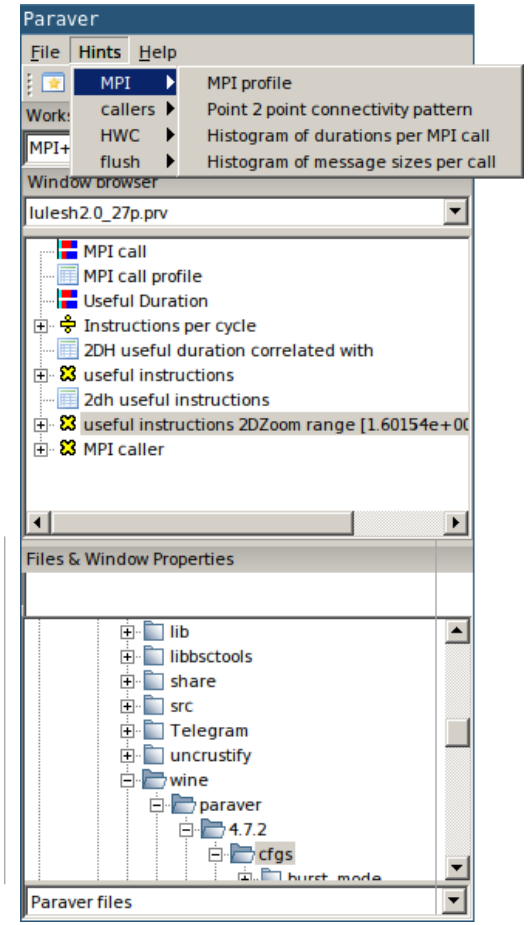

CFG's distribution

- Paraver comes with many more included CFG's